

# CA1 – Image Classification on CIFAR100

Name: Quah Johnnie  
Class: DAAA/FT/2B/o4  
Admin No: 2007476

# Contents (Jupyter Highlights)




<u>Sections</u>	<u>Remarks</u>
<b>1.0) Dataset Preparation &amp; EDA</b>	Distribution, Glance of Dataset, Outliers (Limitation)
<b>2.0) Data Augmentation</b>	Data splitting & Augmentations
<b>3.0) Modelling</b>	Types of model used & Evaluation
<b>4.0) Model Improvement</b>	Comparing Augmentation & Coarse To Fine Hyperparameter Tuning
<b>5.0) Final Evaluation</b>	Test Evaluation & Feature Mapping (1 <sup>st</sup> layer)
<b>Summary</b>	Conclusion & Rooms for possible improvement

# 1.0) CIFAR100 at a glance (EDA)

Current Highest Validation Accuracy (Fine) : NIL  
Current Highest Validation Accuracy (Coarse) : NIL



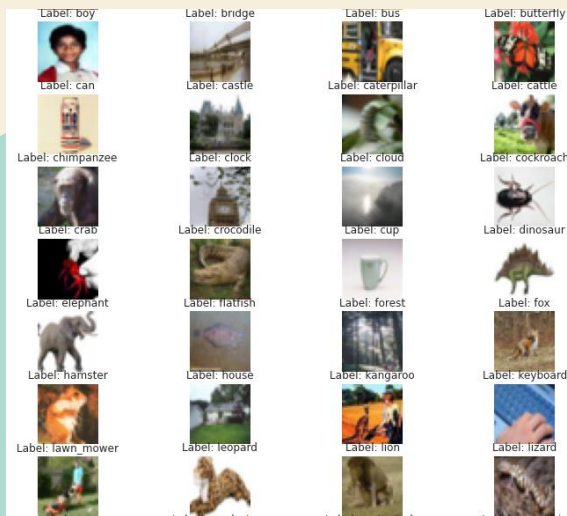
Training data is uniformly distributed



Fine Labelled

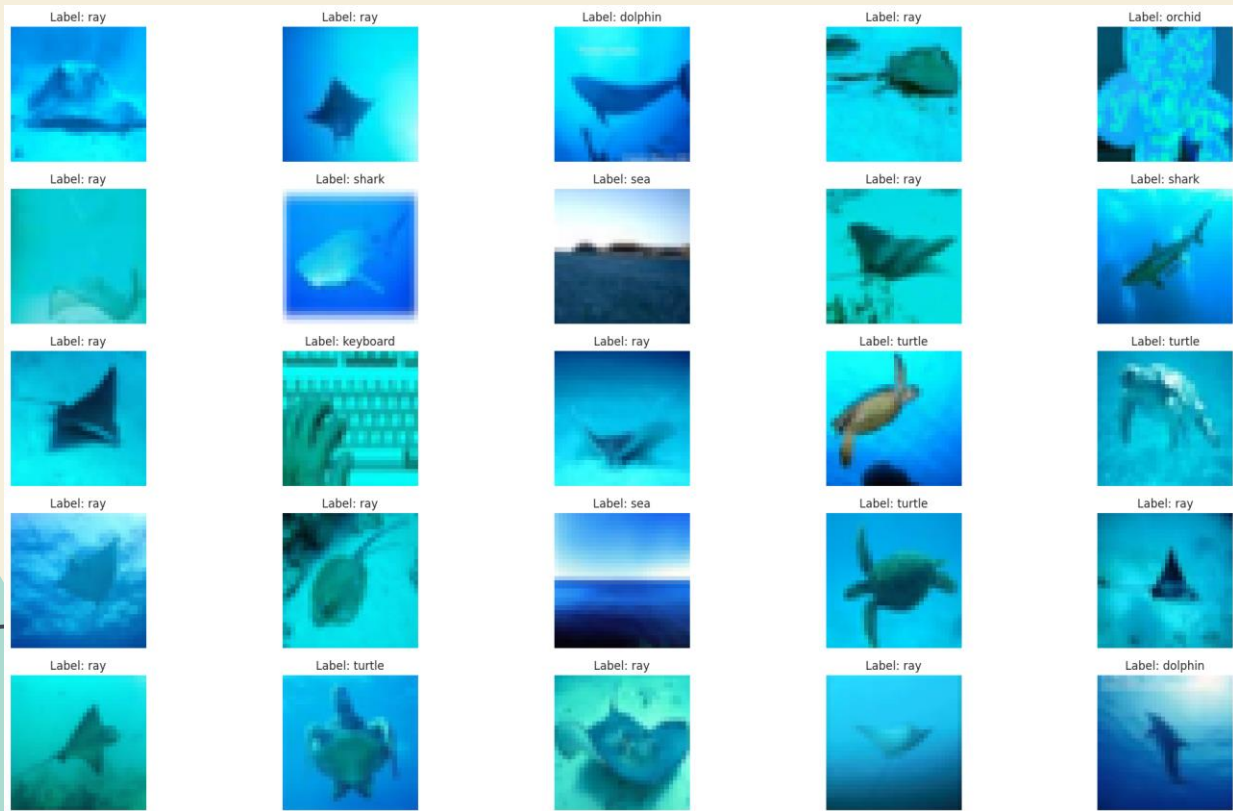
Coarse Labelled

Average of These Labels



# 1.0) Outliers Analysis - Autoencoder

Current Highest Validation Accuracy (Fine) : NIL  
Current Highest Validation Accuracy (Coarse) : NIL



Are these outliers?  
Well yes and no...

How to improve  
outlier analysis and  
look into more fine-  
grained outliers?

## 2.0) Data Augmentation/Splitting

Current Highest Validation Accuracy (Fine) : NIL  
Current Highest Validation Accuracy (Coarse) : NIL

### DATA SPLIT

Training Data split into:

- 40K Training
- 10K Validation
- 10K Testing

### DATA AUGMENTATION

Composing Augmentation Sets With Different Combinations Of:

- ColorJitter
- CutMix
- RandomRotation
- RandomPerspective
- RandomSolarize
- HorizontalFlip
- AutoAugment (CIFAR10)
- AutoAugment (ImageNet)
- AugMix

## 2.0) Visualising Augmentation

Current Highest Validation Accuracy (Fine) : NIL  
Current Highest Validation Accuracy (Coarse) : NIL

Custom: (ColorJitter, RandomRotation, RandomPerspective, RandomSolarize, HorizontalFlip)



AutoAugment (ImageNet Policy)




AutoAugment (CIFAR10 Policy)



AugMix



CutMix



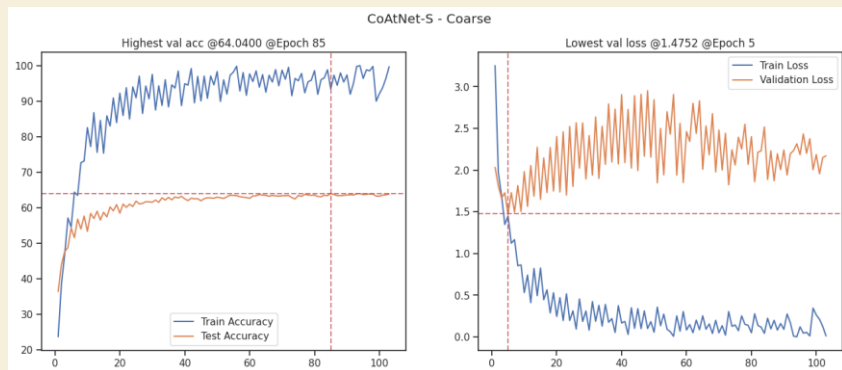
# 3.0) Model & Evaluation

Current Highest Validation Accuracy (Fine) : 54.66%  
Current Highest Validation Accuracy (Coarse) : 68.80%

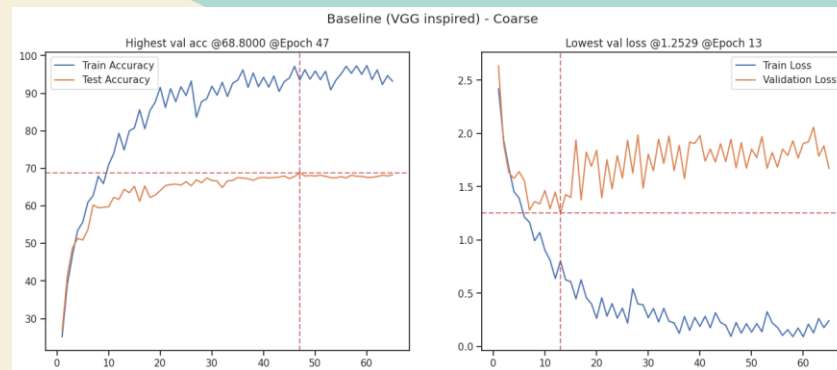
Lowest Val Loss	Highest Val Acc	Lowest Train Loss	Highest Train Acc	Epoch (Highest Val Acc)	Model Description	Parameter
1.981114	53.37	0.095882	<u>97.58750</u>	71	Baseline (VGG inspired)	9,198,276
2.512835	48.86	0.045856	<u>99.10125</u>	54	ResNet34	23,712,932
2.481324	48.89	0.034021	<u>99.35750</u>	49	ResNeXt-M	23,184,804
2.202396	<u>54.66</u>	0.004562	<u>99.94875</u>	80	CoAtNet-S	32,396,096
2.366959	49.75	0.030436	<u>99.52000</u>	72	DenseNet-M	20,013,928
1.252915	<u>68.80</u>	0.092217	<u>97.36000</u>	47	Baseline (VGG inspired) - Coarse	9,198,276
1.560987	60.92	0.026511	<u>99.34250</u>	46	ResNeXt-M - Coarse	23,184,804
1.475237	64.04	0.001255	<u>99.99500</u>	85	CoAtNet-S - Coarse	32,396,096

**\*Extreme overfitting...**  
Perhaps, stronger augmentation has to be done before I see better results on more complex models...

Best coarse labelled model – Learning curve



Best fine labelled model – Learning curve



All models trained with simple augmentation done every 2 epochs.

# 4.0) Model Improvements 1&2

## Comparing different augmentation compositions

Lowest Val Loss	Highest Val Acc	Lowest Train Loss	Epoch (Highest Val Acc)	Model Description	Parameter	Highest Train Acc
2.152134	49.41	0.077599	98	CoAtNet-AutoAugment sample:0	32,396,096	97.837500
2.839644	36.12	0.072459	90	CoAtNet-AugMix sample:0	32,396,096	96.557500
2.238827	46.17	0.626822	68	CoAtNet-ImageNet sample:0	32,396,096	79.947500
2.335603	44.15	0.113020	69	CoAtNet-CustomAugment sample:0	32,396,096	94.617500
1.687326	57.16	0.859597	67	CoAtNet-AutoAugment+CutMix sample:1	32,396,096	83.082500
1.718855	55.71	1.011111	55	CoAtNet-ImageNet+CutMix sample:1	32,396,096	79.366667
1.748327	55.29	0.916276	61	CoAtNet-AutoAugment+ImageNet sample:3	32,396,096	80.396667
1.853471	53.67	0.602639	67	CoAtNet-AutoAugment+ImageNet sample:4	32,396,096	80.353750
1.766168	56.50	0.301244	88	CoAtNet-AutoAugment sample:5	32,396,096	88.788750
1.746079	56.72	0.435966	102	CoAtNet-ImageNet sample:5	32,396,096	85.050000
2.358972	45.45	0.089996	53	CoAtNet-AugMix sample:5	32,396,096	95.651250

Lowest Val Loss	Highest Val Acc	Lowest Train Loss	Epoch (Highest Val Acc)	Model Description	Parameter	Highest Train Acc
0.939657	74.06	0.378762	103	SimpleNet-AutoAugment+CutMix sample:1 (Coarse)	9,198,276	90.644167
0.938005	73.76	0.535824	98	SimpleNet-ImageNet+CutMix sample:1 (Coarse)	9,198,276	88.567500
0.985341	72.69	0.173510	78	SimpleNet-AutoAugment sample:5 (Coarse)	9,198,276	92.812500
1.007861	71.22	0.152900	81	SimpleNet-ImageNet sample:5 (Coarse)	9,198,276	92.907500

## Comparing different complexity of models

Lowest Val Loss	Highest Val Acc	Lowest Train Loss	Epoch (Highest Val Acc)	Model Description	Parameter	Highest Train Acc
1.604467	59.95	0.842774	87	Simpler Custom CoAtNet	-	84.537500
1.470690	63.85	0.607647	123	Simpler Custom CoAtNet2	-	90.815833
1.676687	56.92	0.944950	75	Complex Custom CoAtNet	-	80.855000

Lowest Val Loss	Highest Val Acc	Lowest Train Loss	Epoch (Highest Val Acc)	Model Description	Parameter	Highest Train Acc
0.913700	75.03	0.536745	119	Simpler SimpleNet	-	87.925000
0.918515	74.26	0.592309	64	Complex SimpleNet	-	84.976667

Current Highest Validation Accuracy (Fine) : 63.85% (+9.19%)  
 Current Highest Validation Accuracy (Coarse) : 75.03% (6.23%)

ImageNet Augmentation + CutMix will be used as they provide the lowest variance for both fine and coarse labelled models.

~17M Params

~22M Params

~50M Params

~7M Params

~13M Params



## 4.0) Model Improvements 3&4

### Improvement for fine labelled model

#### Hyperparameter tuning from a checkpointed model:

- The checkpointed model was trained for 200 epochs with a high weight decay
- Checkpointed model reached 63.94% Val Accuracy.
- Augmentation have to be done every epoch during hyperparameter tuning.
- Heavily cut down number of epochs to train.

Trial #32 Finished - Search Time 27.56 Mins  
Total Time Elapsed: 1002.02 Mins

Hyperparameters	Trial Values: #32	Best Trial Values: #20
Learning Rate	0.000316	0.000100
Weight Decay (L2)	0.000000	0.000006
Momentum	0.920000	0.920000
Highest Val Acc	67.26	69.52
Epoch (Highest Val)	40	46

Next trial: [0.00021544346900318823, 4.6415888336127773e-07, 0.9199999999999999]

Trial ended at trial #32

Current Highest Validation  
Accuracy (Fine) : 69.52% (+5.77%)  
Current Highest Validation  
Accuracy (Coarse) : 79.57% (4.54%)

### Improvement for coarse labelled model

#### Transfer learning from our fine labelled model:

- Model is copied from our own hyperparameter tuned fine labelled model.
- Same X\_train image data, different y\_train.
- Change classification layer. Output as 20.
- Val accuracy already hit 75+% during the first 5 epochs
- Hyperparameter tuning can also be done. However, it is not done here due to time limitations. Hence, improvement not as drastic as fine labelled model.

- [Epoch 5/70] | Train Loss: 0.505 | Train Accuracy: 88.33583333333333 | Val Loss: 0.857 | Val Accuracy: 75.79 | Est: 380.75s  
- [Epoch 10/70] | Train Loss: 0.492 | Train Accuracy: 90.48 | Val Loss: 0.764 | Val Accuracy: 78.96 | Est: 371.58s  
- [Epoch 15/70] | Train Loss: 0.475 | Train Accuracy: 90.96333333333334 | Val Loss: 0.729 | Val Accuracy: 79.09 | Est: 358.75s  
- [Epoch 20/70] | Train Loss: 0.455 | Train Accuracy: 91.21666666666667 | Val Loss: 0.730 | Val Accuracy: 79.33 | Est: 357.70s  
- [Epoch 25/70] | Train Loss: 0.464 | Train Accuracy: 90.43166666666667 | Val Loss: 0.722 | Val Accuracy: 79.41 | Est: 364.54s  
- [Epoch 30/70] | Train Loss: 0.388 | Train Accuracy: 92.63166666666666 | Val Loss: 0.730 | Val Accuracy: 79.11 | Est: 358.24s  
- [Epoch 35/70] | Train Loss: 0.435 | Train Accuracy: 91.93916666666667 | Val Loss: 0.741 | Val Accuracy: 79.01 | Est: 344.52s  
- [Epoch 40/70] | Train Loss: 0.499 | Train Accuracy: 89.32333333333334 | Val Loss: 0.726 | Val Accuracy: 79.21 | Est: 365.33s  
EarlyStopper triggered at epochs: 42  
\*No improvement to validation loss and accuracy could be seen for the past 18 epochs  
Highest Val Accuracy: 79.57 @ epoch 24 | Lowest Val Loss: 0.7164515900611877 @ epoch 16

## 5.0) Final Evaluation

All training data is now used for our final model training. Evaluated done once on test data.

The 200 epochs checkpointed model is train with the best hyperparameter for 46 epochs. However, it is better to train from scratch as we have slightly more data for our final training.

### Fine labelled model

```
test_loss, test_accuracy, wrong_samples, wrong_preds, actual_preds, class_dict = eval(finalmodel_fine, criterion, test_loader, val_loader)
print('\nFinal model (fine) test_accuracy (Top-1 accuracy): ', test_accuracy)
print('\nFinal model (fine) test_loss: ', test_loss)
```

Final Model

Final model (fine) test\_accuracy (Top-1 accuracy): 69.64

Final model (fine) test\_loss: 1.30165576338768

Display only top 10 lowest f1 scores

```
displayLowestF1(class_dict)
```

```
otter {'precision': 0.3568181818181818, 'recall': 0.32, 'f1-score': 0.3372340425531915, 'support': 100}
lizard {'precision': 0.40082474226804123, 'recall': 0.39, 'f1-score': 0.3953299492385787, 'support': 100}
seal {'precision': 0.46025641025641024, 'recall': 0.36, 'f1-score': 0.39955056179775285, 'support': 100}
girl {'precision': 0.4657303370786517, 'recall': 0.41, 'f1-score': 0.4315343915343915, 'support': 100}
man {'precision': 0.4575824175824176, 'recall': 0.42, 'f1-score': 0.4379057591623037, 'support': 100}
woman {'precision': 0.44, 'recall': 0.44, 'f1-score': 0.4400000000000001, 'support': 100}
shrew {'precision': 0.43622641509433965, 'recall': 0.46, 'f1-score': 0.4477669902912621, 'support': 100}
mouse {'precision': 0.4456603773584906, 'recall': 0.47, 'f1-score': 0.45747572815533984, 'support': 100}
boy {'precision': 0.4670833333333333, 'recall': 0.45, 'f1-score': 0.4583673469387755, 'support': 100}
bear {'precision': 0.50987951807228917, 'recall': 0.43, 'f1-score': 0.46622950819672134, 'support': 100}
```

Test Accuracy (Fine):

69.64% (+0.12%)

Test Accuracy (Coarse):

79.65% (+0.08%)

Same procedure is done as before. Transfer learning from our fine labelled model, change last layer and trained for 25 epochs.

### Coarse labelled model

Final model (coarse) test\_accuracy (Top-1 accuracy): 79.65

Final model (coarse) test\_loss: 0.7096987730026245

```
displayLowestF1(class_dict2, classes=20)
```

```
b'reptiles' {'precision': 0.644887983706721, 'recall': 0.634, 'f1-score': 0.6393945509586277, 'support': 500}
b'small_mammals' {'precision': 0.6481226053639846, 'recall': 0.674, 'f1-score': 0.6607827788649706, 'support': 500}
b'aquatic_mammals' {'precision': 0.671980198019802, 'recall': 0.678, 'f1-score': 0.6749751243781095, 'support': 500}
b'medium_mammals' {'precision': 0.7826696832579186, 'recall': 0.71, 'f1-score': 0.7399898089171974, 'support': 500}
b'large_carnivores' {'precision': 0.7027586206896552, 'recall': 0.804, 'f1-score': 0.7496296296296296, 'support': 500}
b'non-insect_invertebrates' {'precision': 0.8040909090909091, 'recall': 0.716, 'f1-score': 0.7572340425531914, 'support': 500}
b'large_omnivores_and_herbivores' {'precision': 0.7487072243346007, 'recall': 0.784, 'f1-score': 0.7659064327485379, 'support': 500}
b'household_electrical_devices' {'precision': 0.8704866180048662, 'recall': 0.728, 'f1-score': 0.7922832052689353, 'support': 500}
b'fish' {'precision': 0.8202061855670103, 'recall': 0.798, 'f1-score': 0.8089340101522843, 'support': 500}
b'insects' {'precision': 0.8826905829596412, 'recall': 0.796, 'f1-score': 0.8368710359408033, 'support': 500}
```

# 5.0) Final Evaluation – Error Analysis

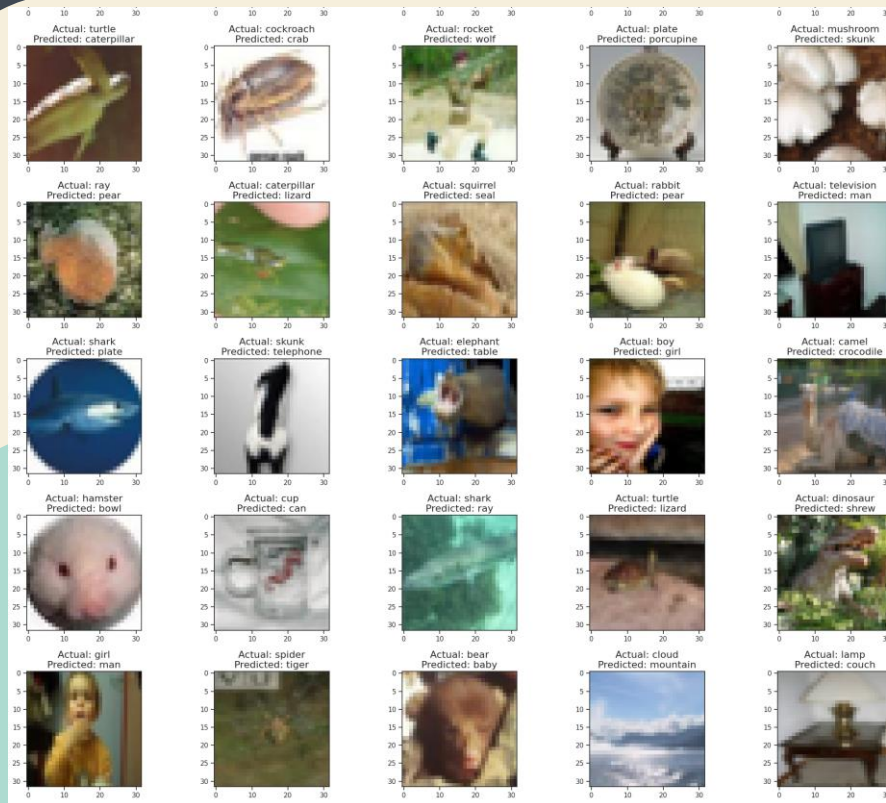
Test Accuracy (Fine):

69.64%

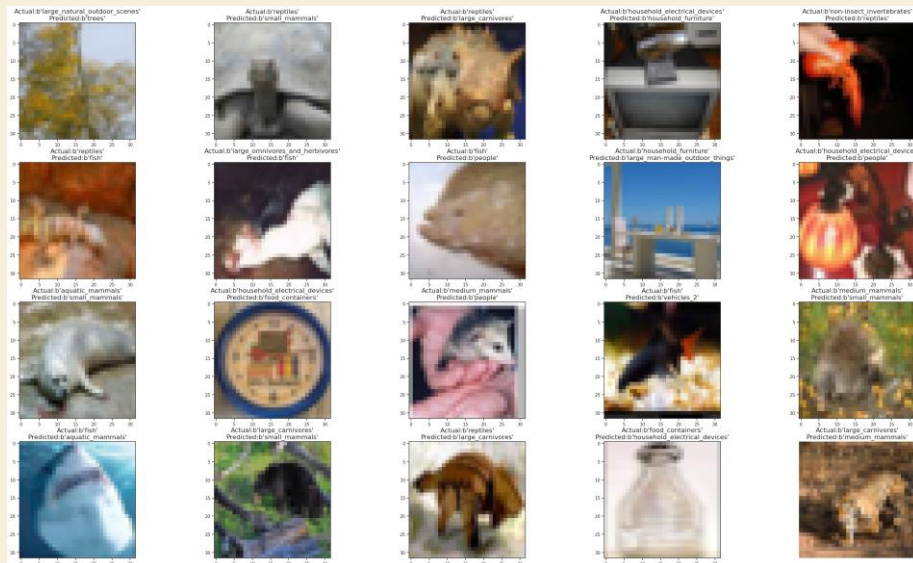
Test Accuracy (Coarse):

79.65%

## Fine label errors



## Coarse label errors



## 5.0) Final Evaluation – Feature Map (1<sup>st</sup> few layers)

Test Accuracy (Fine) :


69.64%

Test Accuracy (Coarse) :


79.65%

Original (3 Channels):

Cattle



3 -> 64 Filters



Batch Normalized +  
Activation Function

## 5.0) Conclusion

### Quick Conclusion:

We did EDA, then splitted our training data into train and validation data and did numerous augmentation compositions. Tried many models before settling on a custom made CoAtNet & VGGNet inspired model. We tested multiple augmentation compositions & also different data replication with augmentation methods, before doing hyperparameter tuning with model checkpoint & transfer learning to improve our model and finally did our final evaluation and analysis of our models.

### Rooms for possible improvement:

Train the model from scratch (instead of checkpoint) since I have the whole training data to work with. Hyperparameter tune coarse labelled model as well. Explore other model scaling methods, most notably EfficientNet compound scaling. Explore more complex augmentation methods, FMix etc. Smaller batch size might help.

\*No matter what. I feel that I have already tried and experiment a lot of ideas during this CA1 but there are still countless ways to improve a model, which is the beauty of deep learning :)

Me: \*uses machine learning\*

Machine: \*learns\*

Me:



*Essential lame memes*



# THANKS!

## Fin. 完。

CREDITS: This presentation template was created by **Slidesgo**, including icons by **Flaticon**, and infographics & images by **Freepik**

Please keep this slide for attribution